

Backend Proposal

Purpose & Scope

Purpose

Provide a lightweight backend service for a web frontend that translates short chat messages (including emoji) bidirectionally between mainstream languages (primarily English). The Slack bot is an optional extension for demonstration only.

Scope

- In Scope:

Provide the following APIs:

Endpoint	Function	Corresponding Feature	Priority & Note
POST: <code>/api/v1/translate/basic</code>	Text2emoji & emoji2text translation	<i>Basic text2emoji / emoji2text</i>	Essential. Core API
POST: <code>/api/v1/feedback</code>	Feedback submission	<i>Feedback & quality + google form / feedback for mode</i>	Essential. Feedback API
POST: <code>/api/v1/translate/contextual</code>	Contextual translation	<i>Added context (previous messages) Generation style selection Country-based contextual translation</i>	Nice-to-have. Provide context for translation, may include regular, style ,user country etc.
POST: <code>/api/v1/models</code>	Get the available model list	<i>A public endpoint to our model</i>	Nice-to-have. Provide model info for developers/users.
(Requires DB) POST: <code>/api/v1/user/history</code>	Get translation history for certain user	<i>Show user history</i>	Maybe. A user system is needed
(Requires DB) POST: <code>/api/v1/user/register</code>	User log in/sign up	<i>Gamify it (achievements/stats)</i>	Maybe. A user system is needed
(Requires DB) POST: <code>/api/v1/user/stats</code>	Get stats and achievements for certain user	<i>Gamify it (achievements/stats)</i>	Maybe. Representing a series of endpoint relevant with gamify features.

- Definition of Done:

All API units passed functional tests with latency under 1000ms.

Architecture / System Design


Architecture Components

- Web Frontend: issues translate requests and displays results. RWD is required. Must also implement the user feedback mechanism and submit data to the API Server.
- Slack Bot Client: Handles all Slack platform-specific interactions (e.g., Slash Commands, message event listening, and context menus). It formats Slack events into standard requests for the API Server and handles the asynchronous response/message update.
- API Server (stateless): exposes a single translation method (or two: Encode and Decode) and a dedicated feedback submission endpoint. It serves as the unified gateway for both Web and Slack clients.
- Translation Orchestrator (inside API Server): builds prompts, concurrently requests external model providers, selects the final output and performs minimal post-processing. Crucially, it is responsible for the 'first-success wins' strategy, including result validation, and standardizing the final structured output.
- Model Provider Adapter: Adapters for multiple lightweight LLM hosts (e.g., openrouter and similar). Each adapter must normalize input requests and parse raw LLM responses to a standardized internal format, shielding the Orchestrator from external API variations.
- Data Persistence Layer (Feedback): A lightweight database (e.g., key-value store or SQLite) dedicated to securely storing user feedback data and related metadata to support ongoing model quality improvement and analysis.
- Optional in-process cache: simple LRU cache for repeated identical requests (not required for MVP).

System Diagrams


High-level Component Diagram

The following diagram displays all the MVP components and their place under the Web and Slack endpoints.



Key design detail (1s responsiveness & stability)

- Broadcast-to-many + first-success strategy: the Orchestrator broadcasts each translation request concurrently to multiple independent lightweight LLM hosts (different providers or instances). The system returns the first successful response that meets minimal acceptance criteria ("first-success wins"). This approach increases likelihood of meeting the ~1s perceived latency target and improves resilience against individual provider slowdown or outage. The following diagram illustrates how Orchestrator process this strategy.



- The frontend expects a synchronous final translation result; the Orchestrator therefore performs parallel external calls but returns only the first acceptable result or a clear timeout/error if none succeed within the allowed window.

API Definition (high-level)

At this point we only designed API definitions for necessary functionalities. We left API nice-to-have and maybe functionalities for the following work.

Core Translation API /api/v1/translate

- Purpose: Unify the text2emoji & emoji2text translation, leave room for future features.

Field	Role	Description
Request Field: originalMessage	Core Input	The text or emoji chain to be translated.
Request Field: isToEmoji	LLM Direction Control	true = Encoding (Text2Emoji); false = Decoding (Emoji2Text).
Request Field: chatHistory	Nice-to-have Placeholder	(Optional) Array of recent messages for Contextual Translation. The system will Limit chatHistory length and truncate aggressively (e.g., last 0-2 messages) to control latency and token usage.
Response Field:	Core Output	The final translated result (plain text or

translatedMessage		emoji chain).
Response Field: metadata.tone	LLM Metadata	The assessed emotional tone (e.g., "Extreme Laughter", "Mild Sarcasm").

Feedback API /api/v1/feedback

- Purpose: Collects data crucial for model quality improvement, fulfilling the **feedback & quality** essential requirement.

Field	Role	Description
Request Field: originalInput	Data Logging	The original text/emoji input that triggered the translation.
Request Field: correctionText	Core Feedback	The correct/suggested translation provided by the user.
Request Field: anonymousId	User Tracking	An anonymous identifier used to track the source of the feedback.
Request Field: rating	User Tracking	An anonymous score provided by user.
Response: status	202 Accepted	Returns a simple success confirmation.

Unified Error and Performance Behavior (Critical for Stability)

- Expected Normal Behavior: Under the Broadcast-to-many Strategy, the system returns the first successful result that passes the Orchestrator's internal validation. The frontend expects a synchronous final result.
- Expected Error Behavior:
 - Bad Request (400): Client input validation failures (e.g., missing fields).
 - Provider Timeout (503): Broadcast Strategy Failure—All external LLM providers have timed out or failed within the allowed time window.
 - Single Provider Failures: Handled internally by the Orchestrator and not exposed to the frontend.

Risks Analysis

- External provider latency and availability variability.

- Mitigation: concurrent requests to multiple providers and timeouts with clear frontend messaging.
 - Avoid multi-agent or multi-step flows in the request path, as they tend to increase latency and complexity.
- Cost and rate limits from third-party APIs.
 - Mitigation: prefer lightweight providers for defaults, cache repeated requests, and apply rate limiting as needed.
 - Translation quality for emoji, slang, and mixed-language inputs.
 - Mitigation: validate with representative examples and iterate on wording/templates.
 - Privacy concerns when sending user text to third parties.
 - Mitigation: document external API usage and avoid storing raw sensitive text where possible.